

End-to-End DevOps CI/CD Pipeline on AWS

Herovire Capstone, Batch 15

Rohit Gupta

What I built

A pipeline that takes my app from a `git push` to a running, monitored service on AWS, without any manual steps.

What it had to do:


- Containerize the app and build it automatically
- Create all AWS infrastructure from code, not the console
- Deploy to a Kubernetes cluster
- Monitor it and alert when something breaks
- Stay cheap and tear down cleanly

Architecture

```
Developer
  | git push
  v
GitHub --> Jenkins (EC2) --> Docker build --> Amazon ECR
                                         |
Terraform (VPC, EKS, IAM, ECR) <-- provisions --|
Ansible (configures Jenkins host)               |
                                         v
                                   kubectl / rollout --> Amazon EKS
                                         |
Prometheus + Grafana <-- scrapes /metrics
```

Herovire DevOps Capstone — End-to-End CI/CD Architecture on AWS

GitHub → Jenkins → Docker → ECR → EKS · provisioned with Terraform · configured with Ansible · monitored with Prometheus + Grafana



Tools I used

Stage	Tool	What it does
Build	Docker + Jenkins	Build the image, push to ECR
Infrastructure	Terraform	VPC, EKS, ECR, IAM, state in S3
Configuration	Ansible	Sets up the Jenkins server
Deploy	kubectl and EKS	Rolls the image onto the cluster
Monitoring	Prometheus + Grafana	Metrics, dashboard, alerts

The app itself is a small Flask service with `/`, `/health` and `/metrics`.

Build and deploy (Jenkins)

The `Jenkinsfile` pipeline runs:

Checkout -> Build -> ECR login -> Tag and Push -> Deploy -> Smoke Test

- Images go to ECR, tagged with the build number
- Deploy runs `kubectl apply` then updates the image tag
- `kubectl rollout status` gates the build, so it only passes if pods come up healthy
- If the smoke test fails, the pipeline rolls back with `kubectl rollout undo`

Infrastructure as code (Terraform)

Everything on AWS is defined in `terraform/`, split into two root modules with separate state:

Module	Contains	Applied by
<code>bootstrap/</code>	VPC, subnets, NAT, Jenkins EC2	me, from a laptop
<code>platform/</code>	EKS, ECR, OIDC role	the Jenkins pipeline

State lives in a versioned S3 bucket, so the environment is reproducible.

`terraform plan` costs nothing, so I only ran `apply` when I was actually working.

The bug that made me split it

My infra pipeline ran `terraform apply` over the **whole** stack, including the Jenkins server it was running on.

I changed the Jenkins `user_data`. Terraform correctly decided the instance had to be replaced, and stopped the machine that was executing the build.

The pipeline destroyed its own executor.

- Not an infinite loop: apply is idempotent, so it is a no-op when the host is unchanged
- It is a self-destruct on *self-modification* - it only appears when you change the CI server's own definition

Fix: `platform/` reads `bootstrap/` but never applies it, so the Jenkins host is not in the state Jenkins touches.

Configuration (Ansible)

Two layers on the Jenkins host:

- **EC2 user data** bootstraps it at first boot: Java 21, Jenkins, Docker, Terraform
- **Ansible** (`configure-nodes.yml`) then installs Docker, kubectl and the AWS CLI, and points kubeconfig at the EKS cluster

Ansible is idempotent, so re-running it converges to the same state. User data only ever runs once, on first boot, which is why the ongoing config lives in Ansible.

Kubernetes

- Deployment with 2 replicas
- Service of type LoadBalancer, so the app is reachable publicly
- HPA to scale on CPU (needs metrics-server installed)
- Liveness and readiness probes on `/health`

Monitoring

- Installed kube-prometheus-stack with Helm (Prometheus, Grafana, Alertmanager)
- The app uses `prometheus-flask-exporter` to expose `/metrics`
- A ServiceMonitor tells Prometheus to scrape it
- Grafana dashboard shows request rate: `rate(flask_http_request_total[1m])`
- Alert `AppPodDown` fires when the scrape target stops responding

Testing

- `tests/smoke_test.sh` checks that `/` and `/health` return 200
- `tests/infra_test.sh` checks nodes are Ready, pods Running, LB is up
- The smoke test runs as a Jenkins stage, so a bad deploy fails the build
- I ran the full cycle end to end: apply, test, destroy

Bonus: GitOps with GitHub Actions and ArgoCD

I also built a second, more modern path and ran it live:

- GitHub Actions builds and pushes the image, authenticating with OIDC so no AWS keys are stored in GitHub
- It then updates the image tag in `k8s/deployment.yaml` and commits it
- ArgoCD watches the repo and syncs the cluster to match

One push took the app from v1 to v2 with no `kubectl` from me.

Demo

1. App running as v1 on its LoadBalancer URL
2. Push a small change
3. GitHub Actions goes green
4. ArgoCD flips from OutOfSync to Synced, pods roll
5. Refresh, app is now v2
6. `kubectl scale --replicas=1`, ArgoCD puts it back to 2

Problems I ran into

- My Mac builds arm64 images but the nodes are amd64, so pods failed to pull.
Fixed by building with `--platform linux/amd64`
- Leftover LoadBalancers blocked the VPC from deleting. I now delete the Services before running destroy
- Prometheus was scraping nothing because the Service was missing an `app` label
- The ArgoCD CRD was too big for a normal apply, needed `--server-side`
- ArgoCD kept recreating things during teardown until I deleted the App first

Cost

- Running cost is roughly \$0.25 to \$0.30 per hour (EKS, 2 nodes, NAT, LB)
- Writing code and running `plan` is free, only `apply` costs anything
- Kept it small: 2 t3.medium nodes, 3 day metrics retention, one NAT gateway
- Destroyed everything after each session, so a full build and test cycle cost about \$0.25
- ECR lifecycle policy clears out untagged images

What I learned

- Infrastructure as code makes environments reproducible and disposable
- Terraform and Ansible solve different problems
- Push based delivery (Jenkins) and pull based delivery (ArgoCD) are both valid
- OIDC and short lived credentials are better than storing secrets
- A deploy is not finished until you can see that it is healthy
- Cost control has to be a habit, not something you fix at the end

Summary

GitHub to Jenkins to Docker to ECR to Terraform to Ansible to EKS, with Prometheus and Grafana on top, plus a working GitOps path.

One `git push` gets me a running, monitored app on Kubernetes.

Thank you. Questions?